Test Report

Generated 2026-10-07 by `python -m tests.make_test_report`. The actual-result column comes from a real run of the suite.

34 test cases, 55 executed checks, 55 passed. `============================== 55 passed in 19.50s ==============================`

Run the suite yourself with:

```
python -m pytest tests -q
```

Scope and approach

No language model is called anywhere in this suite. Agent behaviour is asserted through tool registries and classifier outputs, both deterministic, because a test whose result depends on generation temperature cannot act as a regression gate. Uploads are written to a temporary directory, so no test touches the development database or the real document store.

Two cases (TC-05, TC-14) assert the *presence* of known defects rather than their absence. They document measured problems and fail if those problems grow worse.

Test cases

ID	Category	Input	Expected	Actual
TC-01	Validation	Generated document corpus	4,950 documents, 9 classes, parallel arrays aligned	PASS
TC-02	Validation	All three datasets	Zero empty strings, zero NaN in features or target	PASS
TC-03	Validation	Document labels	Every class within 1% of the mean count	PASS
TC-04	Validation	Document texts	Zero exact duplicates	PASS
TC-05	Known defect	Message texts	Duplicates present but under 10% (known defect)	PASS
TC-06	Validation	Risk feature matrix	Float dtype, all finite, 12 columns, no constant column	PASS
TC-07	Validation	Risk target	Values in {0,1}, positive rate between 20% and 80%	PASS
TC-08	Boundary	Each bounded feature	Values inside their declared domain	PASS (6 runs)
TC-09	Preprocessing	Document texts through the vectoriser	Sparse matrix, one row per document, >95% sparse	PASS
TC-10	Preprocessing	Document texts, min_df 1 vs 2	Pruning reduces the feature count	PASS
TC-11	Preprocessing	Risk features through StandardScaler	Every column mean ~0, sd ~1	PASS

ID	Category	Input	Expected	Actual
TC-12	Split	Documents split 75/25 stratified	Class proportions preserved within 1% in both folds	PASS
TC-13	Split	Document train/test folds	Zero records present in both folds	PASS
TC-14	Known defect	Message train/test folds	Overlap present but under 15% (known defect)	PASS
TC-15	Reproducibility	Risk split run twice with seed 42	Identical folds both times	PASS
TC-16	Functional	Each of 7 demo PDFs + its requirement	The verdict named in the filename	PASS (7 runs)
TC-17	Functional	Expired safety certificate	Rejected, reason names the expiry date	PASS
TC-18	Regression	Passport line holding issue and expiry together	Each date gets its correct role	PASS
TC-19	Regression	Document titled 'CERTIFICATE OF ...'	Identifier taken from the field, not the title	PASS
TC-20	Boundary	Zero-byte upload	Rejected, no crash	PASS
TC-21	Exceptional	HTML content declared as application/pdf	Rejected on magic bytes	PASS
TC-22	Boundary	File one byte over the 10MB cap	Rejected and no partial file left on disk	PASS
TC-23	Exceptional	Employee ids with traversal or illegal characters	Rejected	PASS (5 runs)
TC-24	Exceptional	Crafted relative path '../../../windows/win.ini'	Refused, cannot escape the store	PASS
TC-25	Validation	PDF with no extractable text	Never auto-approved; routed to a human	PASS
TC-26	Performance	One demo PDF, warmed	End-to-end analysis under 5 seconds	PASS
TC-27	Functional	One representative message per intent	Routed to the correct specialist	PASS (5 runs)
TC-28	Boundary	'any update?'	Low confidence, not routed to a specialist	PASS
TC-29	Functional	One message per sentiment class	The matching sentiment label	PASS (3 runs)
TC-30	Security	Message containing six identifier types	Every value stripped before leaving	PASS
TC-31	Access control	Same question from two organisations	Different, organisation-correct policies	PASS
TC-32	Access control	Security question from permitted and non-permitted departments	Restricted document hidden from Sales	PASS

ID	Category	Input	Expected	Actual
TC-33	Architecture	Each specialist's tool registry	Capabilities disjoint and enforced by construction	PASS
TC-34	Performance	50 classifications and 50 retrievals	Mean under 150ms per call	PASS

Coverage by category

Category	Cases
Validation	7
Boundary	4
Functional	4
Preprocessing	3
Exceptional	3
Known defect	2
Split	2
Regression	2
Performance	2
Access control	2
Reproducibility	1
Security	1
Architecture	1

What is not covered

Stated plainly so the gaps are not mistaken for passes:

- **No API-level tests.** Endpoints are exercised manually, not by the suite.
- **No authentication tests**, because there is no authentication to test.
- **No frontend tests.** No vitest or Playwright coverage exists.
- **No agent-response quality tests.** Tool boundaries are asserted; whether an

answer is *good* is not, and would need an LLM-as-judge rubric reported as a trend rather than a gate.

- **No model accuracy gate in CI.** Metric thresholds are not enforced on each run.